

AI OPERATIONS PORTFOLIO · USER MANUAL

AI Career Operations System

A multi-agent operating system that runs a career transition like an operations function.

John White

AI Operations & Optimization

[linkedin.com/in/johnwhiteimse](https://www.linkedin.com/in/johnwhiteimse) · github.com/johnwhite-ai/ai-operations-portfolio

Contents

- 01** Executive Summary
- 02** The Problem
- 03** System Architecture
- 04** The Daily Workflow
- 05** A Worked Example
- 06** The Differentiators
- 07** What This Project Demonstrates

01 Executive Summary

The AI Career Operations System is a complete operating system for running a career transition the way you would run an operations function. It is not a job scanner. It is a five-agent suite, two custom skills, a Notion command center with eight linked databases, a multi-source intake pipeline, a scoring rubric that enforces hard-won personal criteria, and an application-package generator — coordinated into one daily workflow.

It is in production right now: it ran the search that produced this portfolio and the applications behind it.

THE ONE-LINE VERSION

Most people build a job scraper. This operates a function — scan, score, tailor, apply, track, network — with a single source of truth and a daily cadence, designed to protect against the wrong roles, not just surface the most.

02 The Problem

A career transition produces the same failure mode as any unmanaged operation: noise, inconsistency, and no system of record. Hundreds of postings, most wrong on role, industry, compensation, or location. Resumes drift out of sync with the roles they target. Applications get tracked in your head until they don't.

And the real risk underneath all of it: taking another misaligned job because evaluating each opportunity rigorously does not scale by hand. The last bad fit is exactly what this system exists to prevent.

03 System Architecture

Five specialist agents, two skills, a multi-source intake pipeline, and a Notion command center as the single source of truth — coordinated into a daily pipeline. The human supplies judgment; the system supplies synthesis, scoring, tracking, and drafting.



Sources feed the agents and skills, which drive a daily pipeline that syncs to the Notion command center.

The five specialist agents

Agent	Role
Career Strategist	Holds the North Star; challenges drift toward misaligned roles. The judgment layer.
Job Scanner	Multi-source daily scan; filters, scores, writes a ranked report, syncs to Notion.
Job-Fit Scorer	Rates every role across five dimensions for a total out of 50, with hard stops.
Resume Positioning	Tailors the resume per posting across three positioning variants.
Portfolio Coach	Advises on the portfolio work itself.

The two custom skills

Skill	Role
ATS Optimizer	Scores a resume against a posting for keyword/format match; targets 80%+; flags real gaps without stuffing.
Resume Bullet Writer	Enforces the bullet standard: strong verb + specific method + quantified outcome.

The command center (Notion, eight linked databases)

Database	Purpose
Target Companies	Watchlist, tiered, with North Star fit and remote status
Job Applications	The pipeline: each role with its /50 score, salary, status
30-Day Roadmap	Milestone tracker for the transition
Scan Log	History of every scan run — trends over time
Resume Tracker	Every version, ATS score, file location, format
Weekly Review	Application metrics + recalibration discipline
Interview Prep	Per-company prep, linked to applications
Networking	Outreach tracking, linked to companies

04 The Daily Workflow

Run as a managed pipeline, not a series of one-off tasks:

```

scan (multi-source) -> hard-filter -> score (/50) -> ranked report -> sync to Notion
                                     |
                               tailor resume -> ATS-optimize -> generate application package
                                     |
                               apply -> log in Notion -> networking outreach

```

Only roles scoring 28+ surface, sorted Pursue Strong / Pursue / Monitor, each with per-dimension scores, why it fits, honest gaps, and a recommended next action. The scan syncs to Notion automatically every run.

05 A Worked Example

A LinkedIn alert email arrives for "AI Operations Lead at [Company]." The email has a company and a title but no job description.

- **Job Scanner** extracts the (company, title) pair from the Gmail message.
- Because the LinkedIn link is behind a login wall, the scanner finds the real public posting on the company's ATS instead.
- It fetches the full description and runs the hard filter — not quota-carrying, not below the comp floor, right industry. It passes.
- **Job-Fit Scorer** rates it: Role Fit 9, Background 9, Industry 8, Compensation 10, Location 9 — 45/50, Pursue Strong.
- The role tops the ranked report and syncs to the Notion Job Applications database (status: Drafting).
- The human runs Resume Positioning + ATS Optimizer, generates the package, applies, and flips the status to Applied.

THE POINT

A wall — LinkedIn blocking automated access — routed around with a different source. The role was surfaced, scored, tracked, and drafted, with the human only making the apply / no-apply call.

06 The Differentiators

DIFFERENTIATOR 1 — A FULL OPERATING SYSTEM, NOT A TOOL

Five agents, two skills, eight databases, a daily cadence — coordinated. Most people build a scraper. This is the difference between automating a task and operating a function.

DIFFERENTIATOR 2 — IT ENCODES JUDGMENT, NOT KEYWORDS

The scoring rubric carries hard-won personal criteria, including an explicit hard-stop for the kind of role that was a bad fit last time. Its job is to protect against the wrong roles, every run — not to find the most.

DIFFERENTIATOR 3 — IT CLOSES EVERY LOOP

Scan results, application status, resume versions, networking, and weekly metrics all live in one linked system of record. Nothing is tracked in someone's head.

DIFFERENTIATOR 4 — PROVEN BY ITS OWN OUTPUT

This system ran the search that generated this portfolio and the applications behind it. It is in production on the most important project its author has.

07

What This Project Demonstrates

Capability	How this project proves it
Multi-agent system design	Five coordinated specialist agents with a judgment layer
Operations thinking	A job search run as a managed pipeline with a system of record
Encoding judgment into a system	The North Star rubric enforces personal criteria automatically
Multi-source integration	ATS boards, Indeed, Gmail/LinkedIn, and Notion (read + write)
Working around constraints	LinkedIn blocks bots, so the system finds the public posting instead
End-to-end ownership	From scan all the way to a ready-to-submit application package